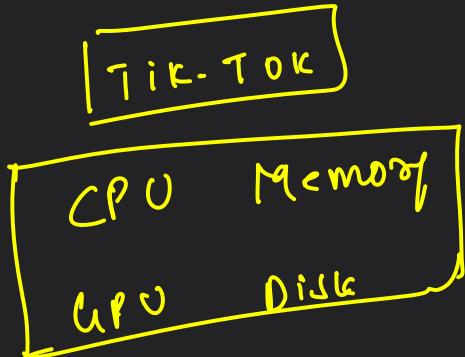


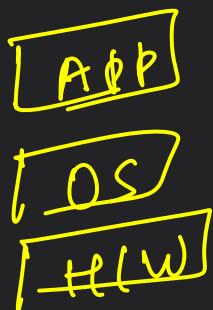
Operating Systems

P. Sailesh Kumar

Operating Systems

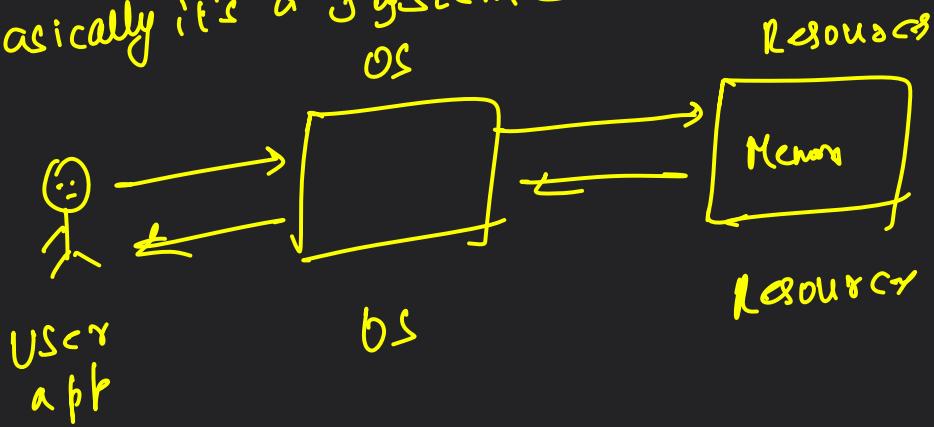


Tik-Tok App
can directly access
the resources(HW) directly.



The functionality of OS
is resource mgmt

* OS is an interface b/w App & HW
(Basically it's a system software)



what happens when OS isn't there?

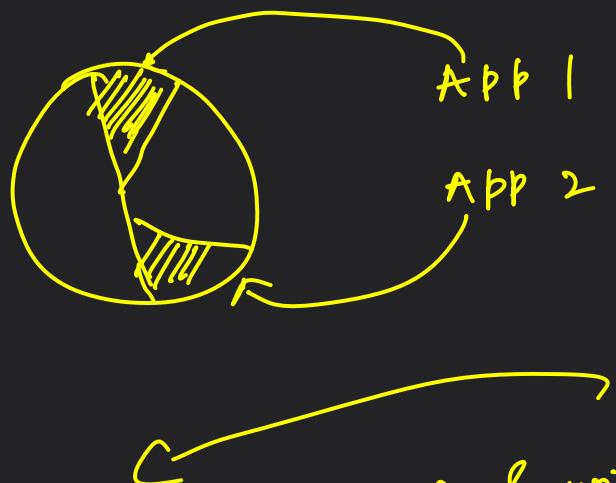
App 1 App 2
↳ write memory myvar
↳ write resource myres1 ↳ write resource myres1
↳ write resource myres2 ↳ write resource myres2
Every developer will write separate code

So this app becomes more bulky.

DRY (Do not repeat yourself)

* To solve this all kinds of memory, obscure mgmt is done by the OS.

Isolation & Protection:



with OS, we can have different memory segments so that there is isolation or else it can write into other app's by trespassing memory.

All this can lead to security loop holes.

* OS provides the abstraction, hiding the underlying HW.

* Basically OS make the work easy for app's developer

Types of OS

1) General purpose OS: Windows, Linux distributions, macOS, iOS, Android

2) Embedded OS: RTOS, QNX

Multi-programming OS

Allows multiple programs to run in 1 CPU
(Serial, pseudo parallel)

Multi-processing OS

Multi-CORE
(more than 1 core)
(Real parallelism)

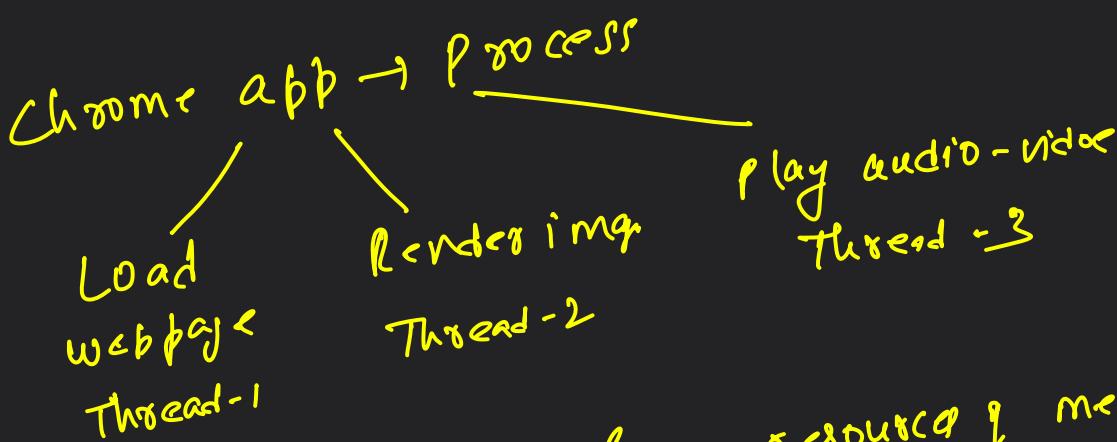
Multi-Threading

process??

- ↳ Basically a program under execution is called as process.
- ↳ Probably running an a.out is called a process

Thread?

- ↳ lightweight process
- ↳ part of process sub-process also called as thread



These threads share same resource of memory, but they run independently.

- * Multi-threading is efficient in multi-core CPU's
- * Multi-threading takes time in single core CPU's
- Now how does OS scheduler decide which thread gets the CPU?

It depends on Scheduling algorithms

Eg process A → Threads A1, A2

process B → Threads B1

process C → Threads C1, C2, C3

Total 6 threads

4 - cores (CPU)

core 0 → A1

core 1 → B1

core 2 → C1

core 3 → C2



core 0 → A2

core 1 → C3

core 2 → A1 again

core 3 → D1 again

These are done by Scheduling algorithms

Multi-processing refers to running multiple process at a time

Eg Running Spotify, same time browsing in chrome, and also reading in zoom call
(P1, P2, P3)



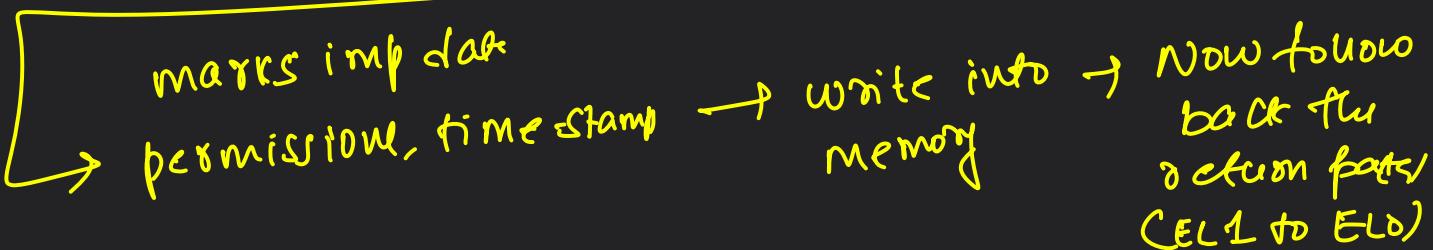
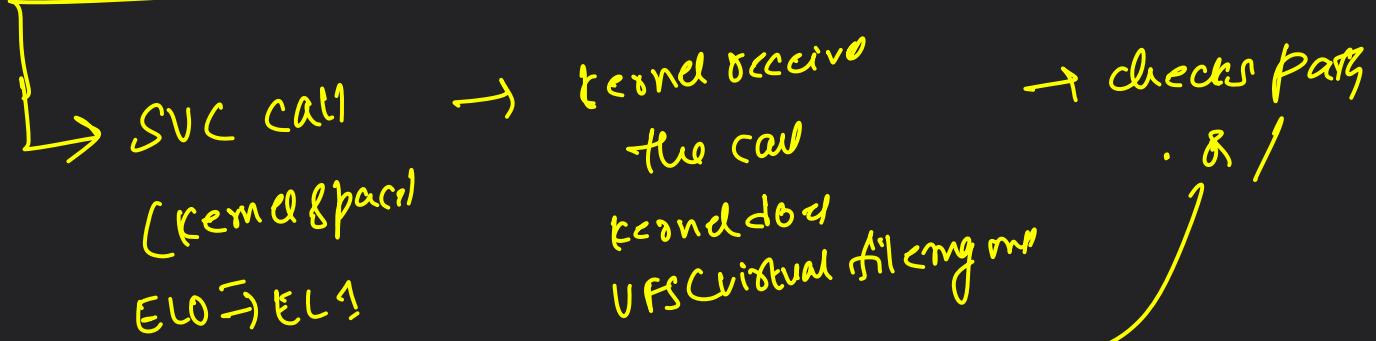
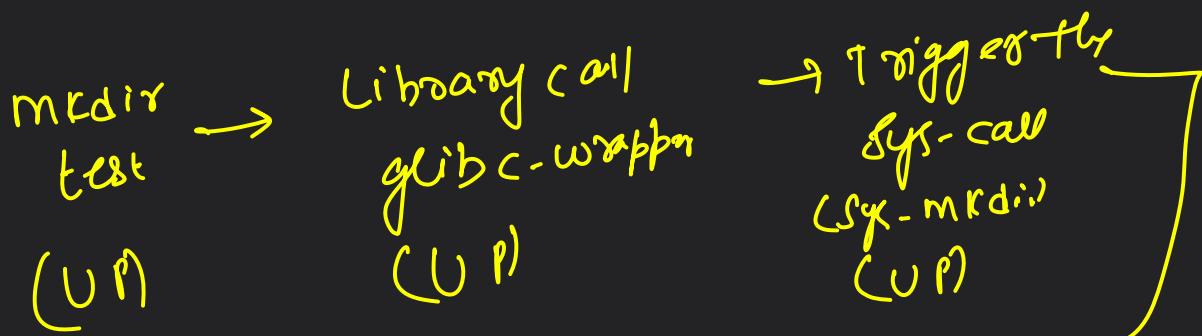
Components of OS

① User Space → User app's

② Kernel Space → Has the access to underlying HW.

User Space: → NO HW Accd

↳ Environment for app's (like GUI / CLI → command line interface)



0 get Val + Success

if fail a already exist → -1

functionalities of Kernel

- 1) Process mgmt (Process creation, termination, comm'n, Process scheduling)
- 2) Memory mgmt : (Allocate, Deallocate, memory free)
- 3) File mgmt: create, delete files
→ Directory mgmt
- 4) I/O mgmt: USB plug in, Game-Controller, Pen-drive
All the I/O mgmt done through Kernel

Types of Kernel

Monolithic Kernel

Entire OS runs on Kernel

More bulky

Eg: Server, Cloud, IoT
Android

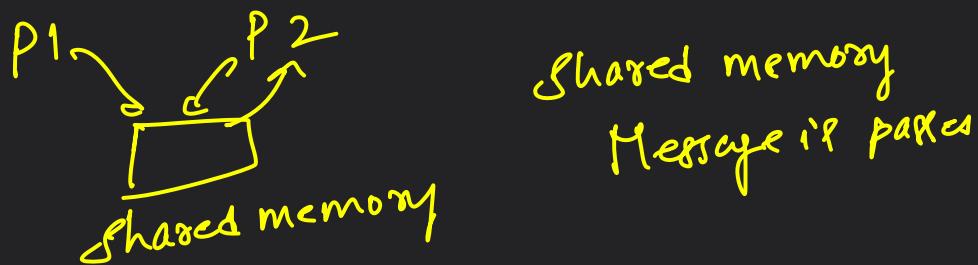
Micro-Kernel

only core fun
like Memory
Process mgmt is
done in kernel
is open

(Less bulky)

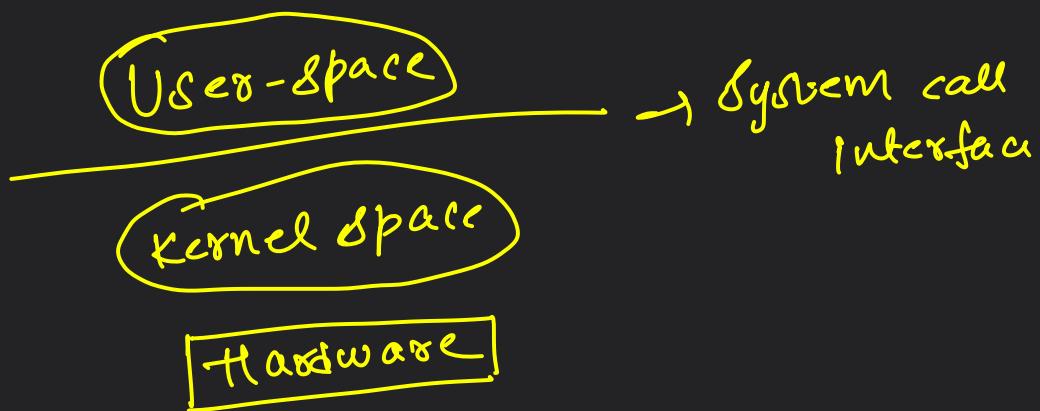
Eg: Multiple context
switching, overhead

IPC (Inter process comm)



i) Message passing through pipe

System-calls

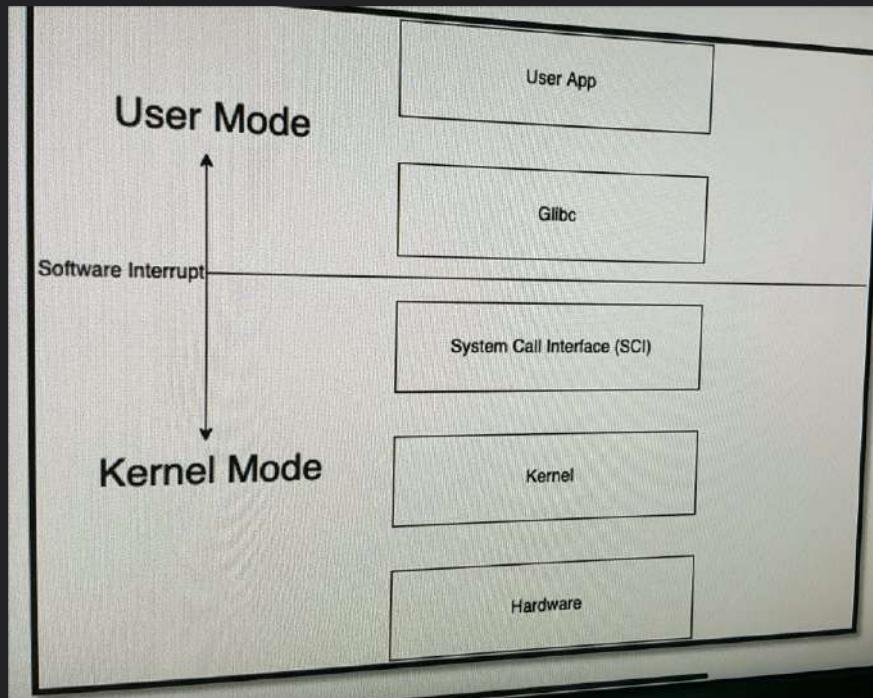


Actual implementation of sys-calls are implemented in 'C'-Language.

Eg . / a.out → Sys-call (SCE) → Kernel mode → Process mgmt
 ↴ Create process → KM to UM
 (Software interrupt/Erip)

SVC in ARM

Sys-call in x86



	Windows	Unix
Process Control	CreateProcess() ExitProcess() WaitForSingleObject()	fork() exit() wait()
File Management	CreateFile() ReadFile() WriteFile() CloseHandle() SetFileSecurity() InitializeSecurityDescriptor() SetSecurityDescriptorGroup()	open() read() write() close() chmod() umask() chown()
Device Management	SetConsoleMode() ReadConsole() WriteConsole()	ioctl() read() write()

How OS boots up

5 major steps

① Power ON → power supply
(Motherboard
Storage
devices)

② CPU loads BIOS/UEFI
(A small program)
→ stored in BIOS chip
(Non-volatile chip)

(BIOS
Basic I/O
UEFI Unified
extensible firmware)

CPU init goes to BIOS

③ BIOS/UEFI runs the code in the flow

3.1 CMOS battery

3.2 POST → Power on Self Test

④ control move to bootloader

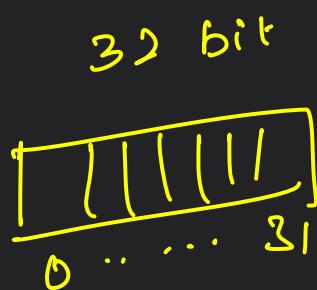
⑤ MBR → Master boot Record
(Disk 0th index)

⑥ EFI → Partition

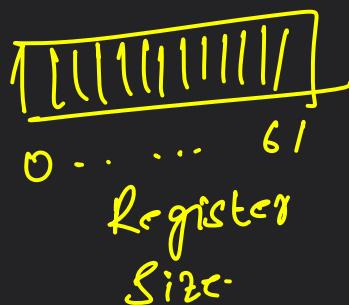
⑦ load the complete OS

32-bit vs 64-bit

64 bit



Register
size



Register
size

2^{32} Unique Address \Rightarrow 4 GB RAM

2^{64} Unique Address \Rightarrow 16 exabytes RAM

why can't we run 64 bit OS in 32 bit Systems

Registers are different CRAX, LBO, R2R on x86-64
E AX, E BX, E SP

Memory addressing - different.

Different Storage in Computer

Registers → Most closest to CPU
(Array kind of structure)
(Instructions are performed)

Cache memory → Additional memory
(for temporary storage to be executed)

- ↳ very fast memory
- ↳ copies frequently used data & instructions.
- ↳ goal is to reduce average time to access memory.

RAM → volatile
↳ used to store the running programs

ROM (HDD / SSD) → long term storage

Virtual memory → part of secondary storage
(HDD / SSD) to extend the RAM.

ROM → Boot ROM (Inside SOC)
firmware

RAM → can be external DRAM

Storage → eMMC (NAND / NOR Flash)
↓
OS, Kernel, App's reside

Process Management

I have an executable (a.out)

process creation

- ① Load the program & static data to memory
 ↳ (RAM)
 used for initialisation

static data

char *p = "Sailash"; static
 int a = 0; Init data

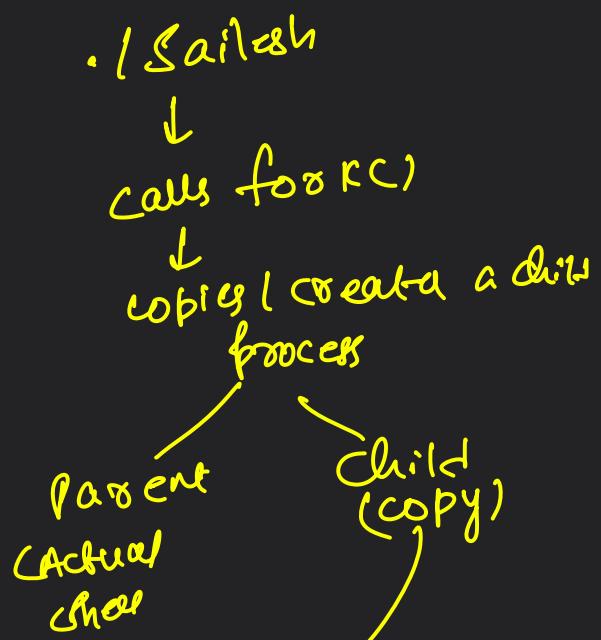
- ② Allocate run-time stack.
 (Local variables, fn args, return values)

- ③ Allocate heap
 ↳ Dynamic allocation

- ④ I/O tasks ↳ C(P), O(P), error
 ↳ handle handle handle

- ⑤ OS hand off control to main

Sailash binary



In child
 Shell calls execv
 ↓
 Kernel reads the ELF
 ↓
 set up text, code, data
 ↓
 Kernel sets up the PC to the = start
 ↓
 control goes back to user & pos
 ↓
 execute code
 ↓
 return 0 (terminate)

fork

Parent
(Actual
shell)

child

(Duplicate of shell)

① parent waits for child to finish using wait system call.

* ② The child is not supposed to remain in memory.
 The child's job is to replace itself with the program we typed

③ Child calls execv, replaces the child memory with our program

④ Kernel looks at the ELF header of the binary to understand
($-start$)

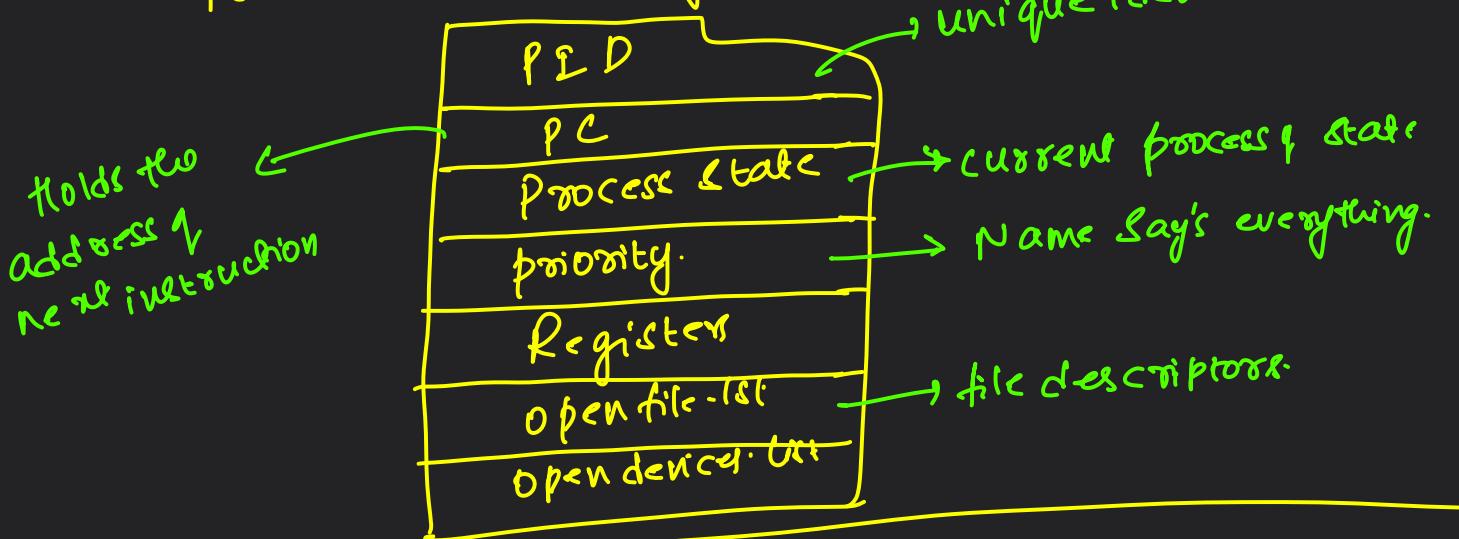
⑤ Set up code/text, data, bss sections.

⑥ PC points to $-start$

⑦ After start, jumps to main

Process control block has all the process attributes.

it is a "struct" type of data structure.



Process-states

Life-cycle of process

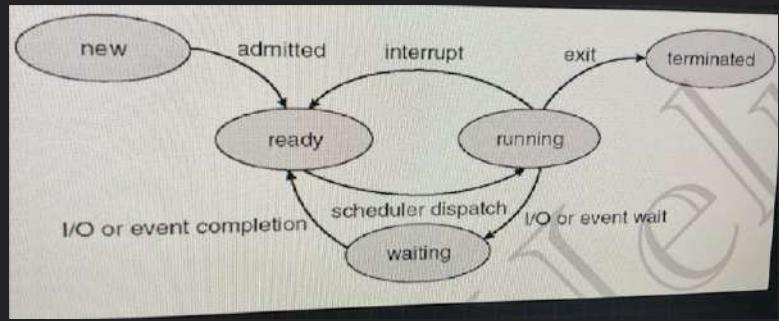
1) New state: Program converting to process

2) Ready state: Process in the memory
(Ready Queue)

3) Running state: P. \rightarrow CPU is allocated

4) Waiting state: Waiting for I/O completion

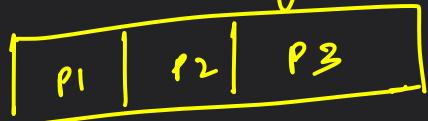
5) Terminated state: Process is finished



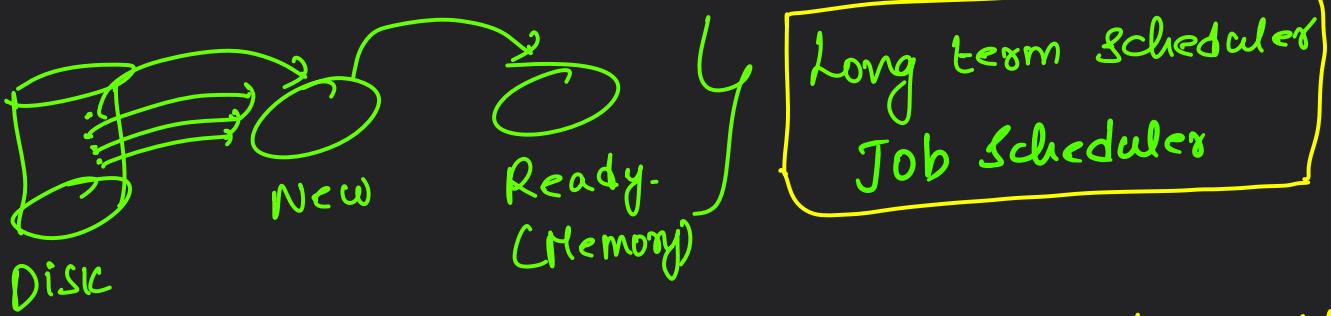
* Running to waiting
I/O

* Running to ready
Time slicing
Again sent back
to memory/Ready Queue

New state process are
in job queue



So if any process goes to ready state
i.e. done through job scheduler



* Now the process is in memory in order to be executed
(Ready Queue)

if should be in Running state (CPU Execution)

Moving from Ready Queue to Running
is done by
CPU Scheduler.

X

Secondary storage

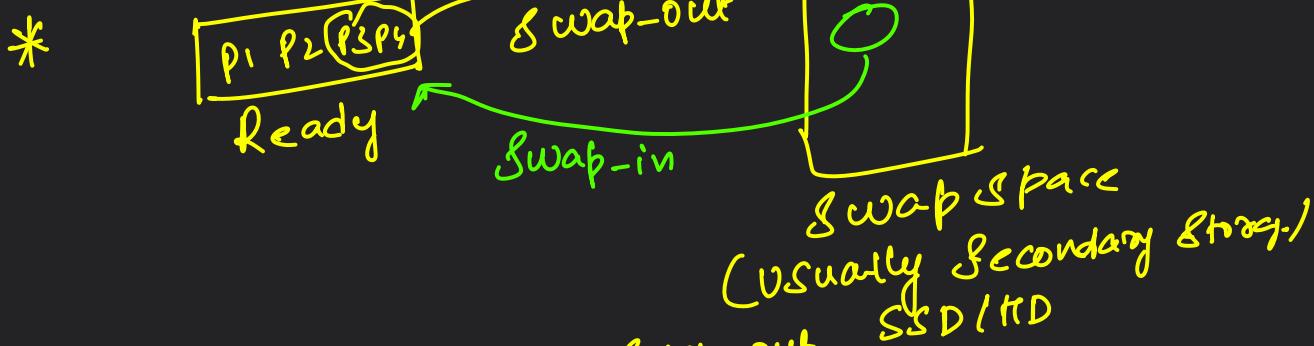
Context-switching



LTS → Mix of processes
(I/O intensive, CPU intensive)

Medium Term Scheduler (MTS)

- * Degree of multi-programming refers to No. of processes in ready queue
- * Now a condition arises that there are few processes in memory which are taking lot of memory. We can't store all of them right? So we need to swap out few of them.

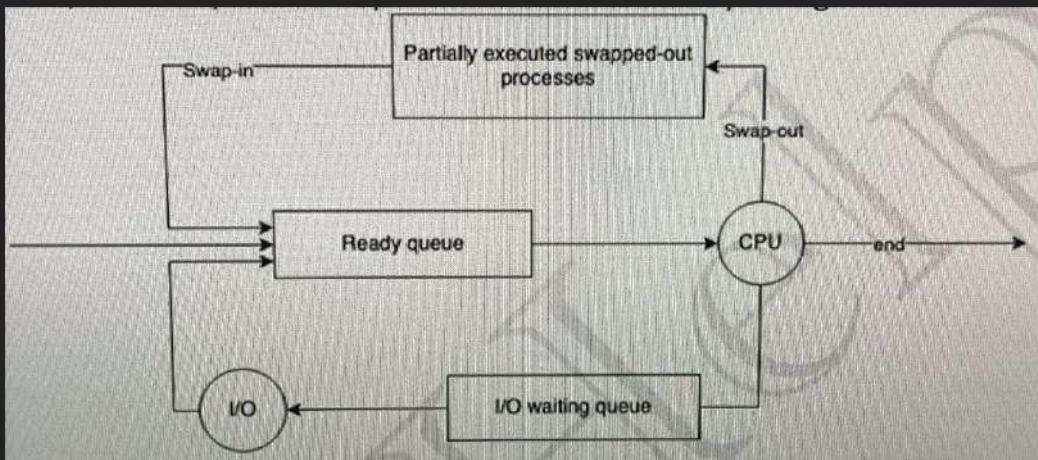


- * We have understood swap-out,

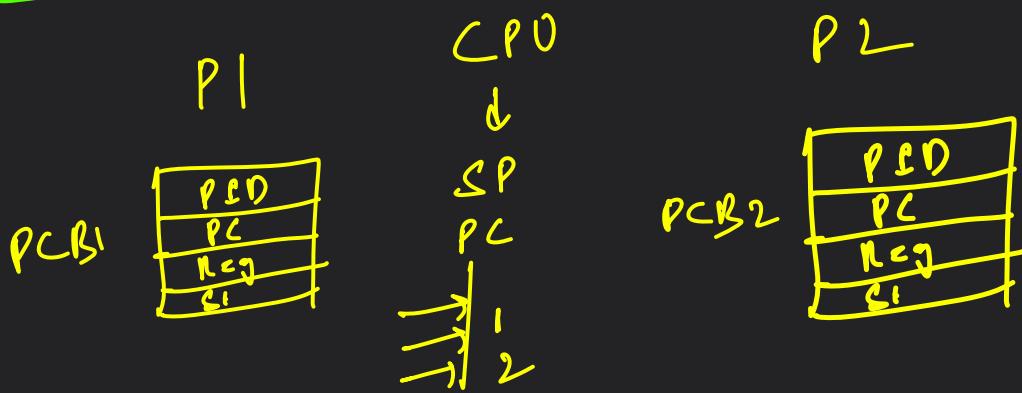
Let's see swap-in

Assume either P1 or P2 is terminated. Now ready queue has memory at that time P3(P4) coming back to ready queue (swap-in).

The entire swap-out to swap-in
called as
Swapping.



Context-Switching



In Detail Analysis of Context Switching

- * A is running in CPU
- * B is Ready (probably created through fork)
+ exec
 ↳ since Ready & created it has its own PCB.
- * Now 'A' is doing some read() operation from disk
(Sys-call)
- * Sys-call is triggered → CPU goes to Kernel mode.

* Kernel Understand's A will take time, so it can't make the CPU idle.

* So Kernel Scheduler will schedule "B".

* Before this it has to store 'A' CPU state.

PC
SP
APR's
CPSR / PSTATE

* Now "A" is mapped as block.

* Scheduled → Preempt → finds → B is given to ready Queue B CPU

Kernel doesn't have any previous context of B
So it loads its initial/primary context

* Now Kernel hits exit

CPU → PC = -8bad1D

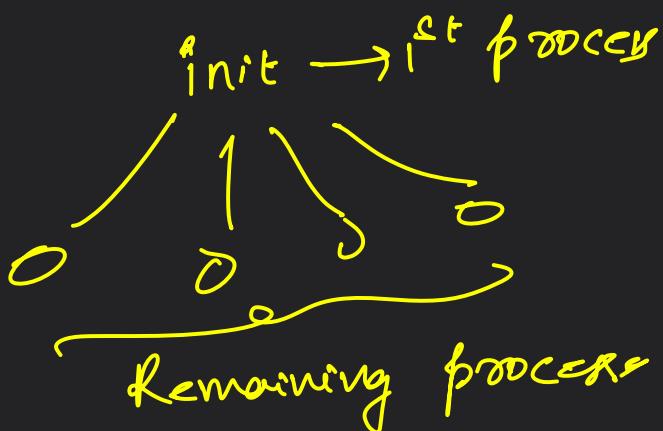
Switch from Kernel to User mode

* Now B starts running on CPU

Disadvantage → overhead (time-waste)
during context switching.

Orphan process:

P1 Running to disk
P2 Running child



→ P1 has suddenly hit an intercept & it has terminated abruptly.

- * Now P1's process has been linked to init.
Orphan process is a process whose parent process has been terminated but child process is still running.

main()

2 pid = fork();

if (pid > 0)

↳ pf("parent existing");
exec(0);

↳ else if (pid == 0)

↳ sleep(10); // give time for parent to exit
pf("I am child, my parent is terminated, my new pid is %d", getpid());

4

Zombie-process

Suppose parent is waiting for 5 min
but child is terminated in memory only

(P1) Parent

fork()

(P2) Child

Child

Executing

Zombie

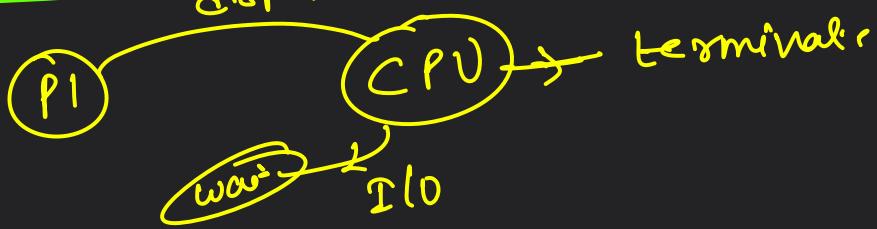
exit

Process Scheduling



which process to pick is done by CPU Scheduler

1) Non-Preemptive:



Time slicing isn't there

2) Preemptive Scheduling: Time Quantizing is followed

Starvation: CPU is not getting the process

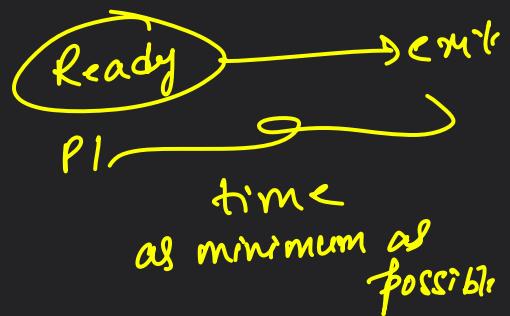
In non-preemptive starvation is more.

In preemptive: Starvation is minimum

CPU Utilization is high, overhead is also more.

* Goals of CPU Scheduling

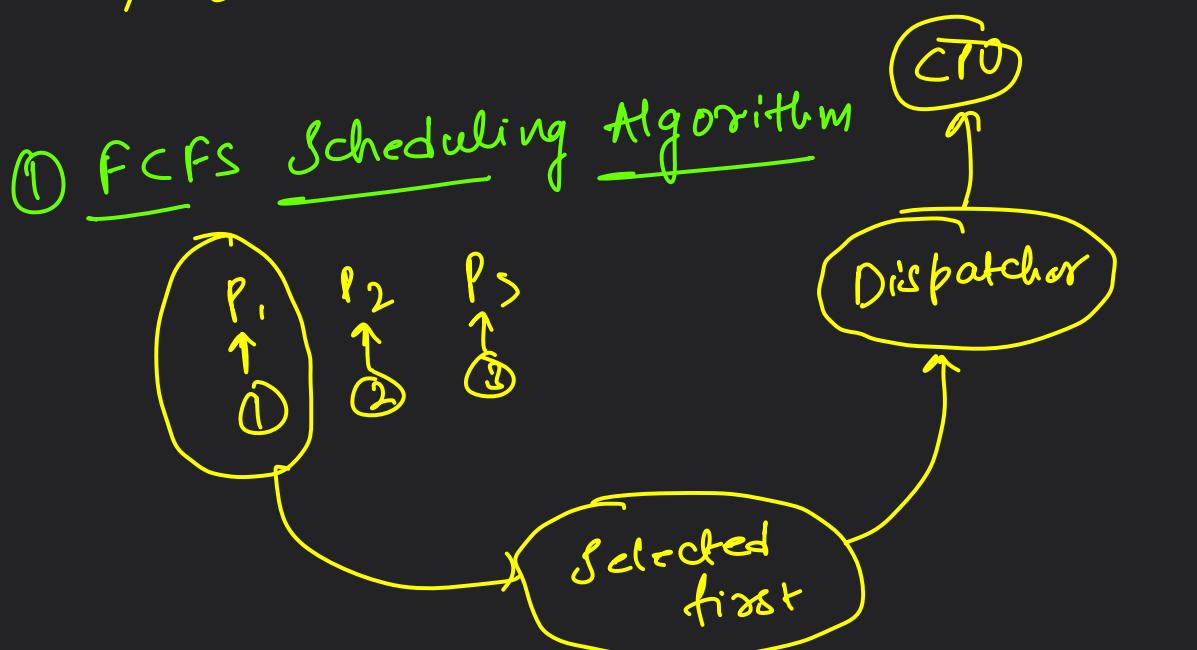
- 1) Max CPU Utilization
- 2) Minimum turn around time (TAT)
- 3) Minimum wait time (WT)



2) Minimum response time (Time b/w Ready queue & CPU getting assigned)

5) Max Throughput: No of processes completed per Unit time

b) Burst-time: Time required for execution.



Pno	AT	BT	CT	TAT	WT
1	0	20			
2	1	2			
3	2	2			

FCFS → Non preemptive

first Arrived goes to Queue CPU later it
like a Queue in ticket counter

SJF (Shortest job first)

process with shortest burst time executes first

Burst time should be known.

Starvation for long jobs.

Round Robin → fixed time Quantum
Issue overhead

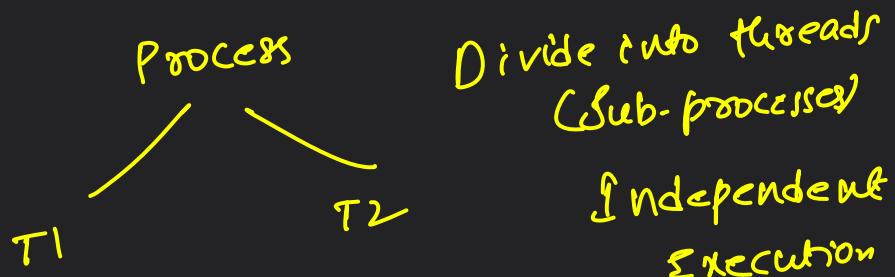
Priority Scheduling:

CPU allocated to the process to higher priority -

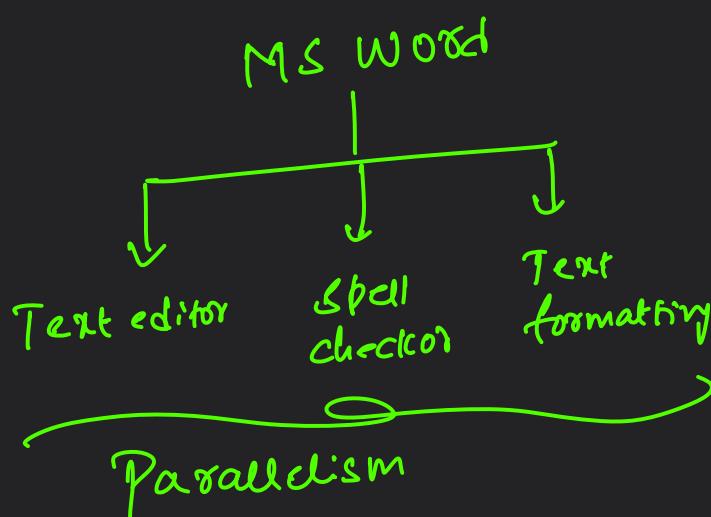
Scheduling in Detail will
be covered later

Concurrency

Ability to execute multiple instructions at a time



why are we dividing into threads?
so that we can divide the task &
try to achieve parallelism.



if not divided
it would
execute
sequentially.

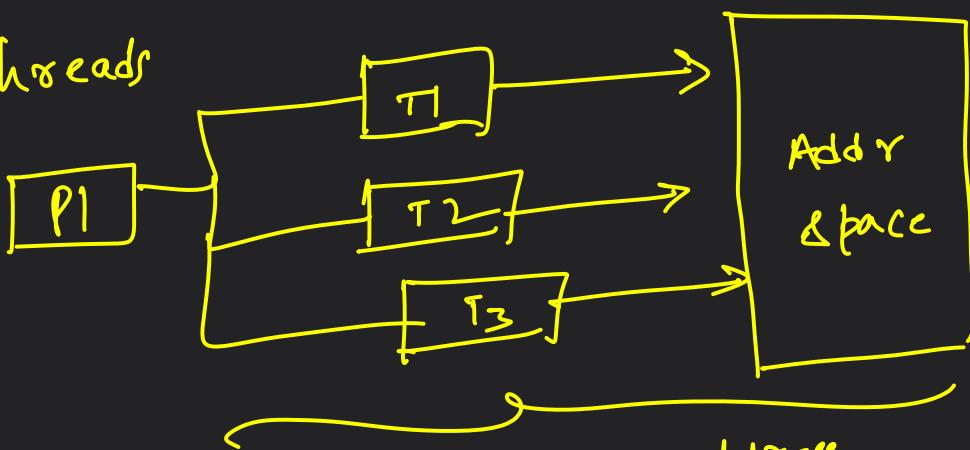
Memory mapping in Process & threads



Mapping is done through
several steps

- 1) compiler/linker → Decides layout of code, data, BSS
- 2) Loader → Loads code & data into memory when process starts
- 3) OS Memory Manager → Provides virtual memory pages, manages stack/heap dynamically.
- 4) MMU → Maps virtual addresses to physical addresses

Threads



using same address
"shared memory"

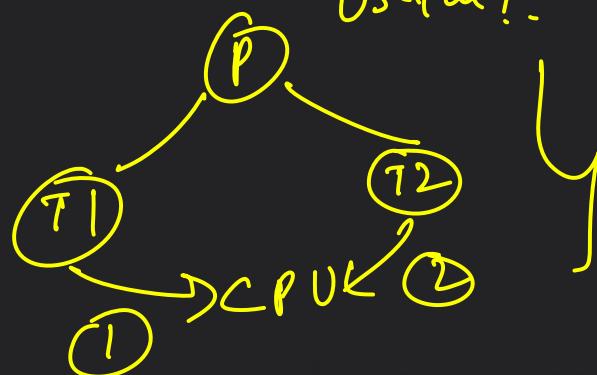
Just like processes,

threads have their own Data structure
"Thread Control Block"

TCB

In thread's context switching only the executing state
(Registers + SP) Address Space remains same

Single CPU → Multithreading
Useful??



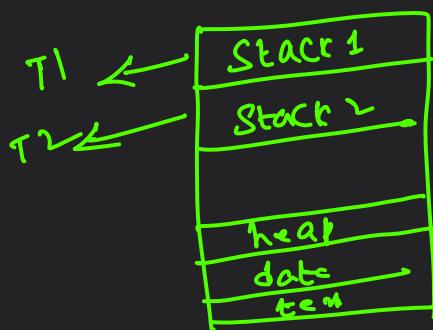
so Basically used
in case of single
core CPU's

Multi-threading is benefit only for multi-core
CPUs

Process



Thread



No of stacks
depend on the
no. of threads

Advantages of Multi-threading

① Responsiveness & Interactive.

② Resource sharing → address space
is shared
(No common gap &
minimal overhead)

③ Thread Economy (context switching has
less overhead)

④ Threads better utilization of multi-core CPU



Gives it to multi-
core
CPU

C → Does not contain built-in support for multi-threading

It entirely relies on the OS

POSIX threads, Pthreads API's are available
Linux, Mac OS Systems.

//

Library is #include <pthread.h>

```
void * task_master(void * arg)
{ pt("Inside task master");
  return NULL;
```

main()

```
# pthread_t tid;
```

```
pt("Inside main before creation");
pt("In main before creation");
```

```
pthread_create(&t_id, NULL, task_master, NULL);
```

```
pthread_join(tid, NULL);
```

```
pt("In main after join");
```

y

* Thread begins task-master

* main doesn't know whether completed or not

* Because task-master is not called by main, OS Scheduler starts new thread that eventually calls the task-master independently.

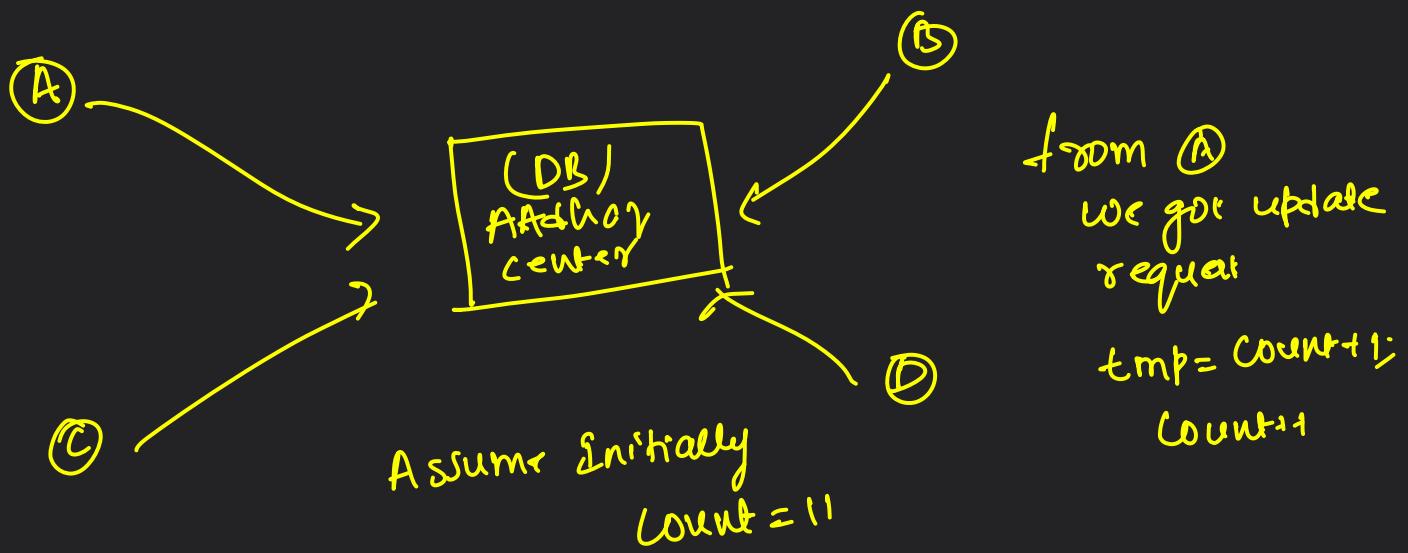
* Main & task-master are running independently.
(Not the expected caller-callee function)

- * So main move to next line naturally.
 - * Here we see pthread-join (This lets main explicitly wait for the thread)
- Synchronisation, OS clean up
Everything is done.

Critical Section

Eg `Count++;` \Rightarrow `Count = Count + 1;`
How this works in CPU-level?

internally
$$\boxed{\begin{array}{l} \text{tmp} = \text{Count} + 1; \\ \text{Count} = \text{tmp}; \end{array}}$$
 Y Internally works like this



from (A)
 we got update request
 $\text{tmp} = \text{Count} + 1;$
 $\text{Count} + 1$

so we are updating tmp to 11 instead of count.

Now B is also updating request
 so instead of count, the tmp will be

$$\begin{aligned} \text{tmp} &= \text{Count} + 1 \\ \text{tmp} &= 12 \end{aligned}$$

(A) ↓
 $\text{tmp } 12$

(B) ↓
 $\text{tmp } 12$

Even after 2 requests our final count is 12, instead of 18

This is race-condition

here DB is critical section

(Multiple threads working on same section/shared memory)

int counter=0;
global var.

void * incr (void * arg)
{
 for (int i = 0; i < 500; i++)
 count++;
}

y pthread_t t1, t2;

pthread_create(&t1, NULL, incr, NULL);

pthread_create(&t2, NULL, incr, NULL);

write join v pthread_join, count;

T1
tmp = count; 0

tmp++; 1

count = tmp;

①

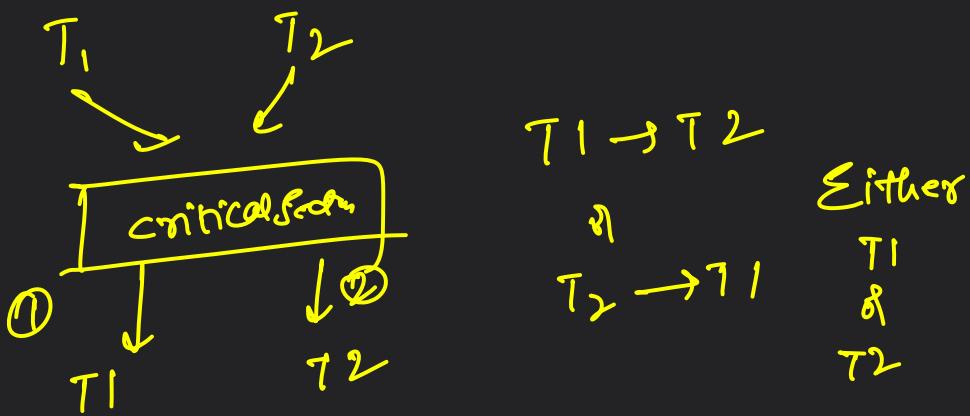
T2
tmp = count; 0

tmp++; 1

count = tmp;

②

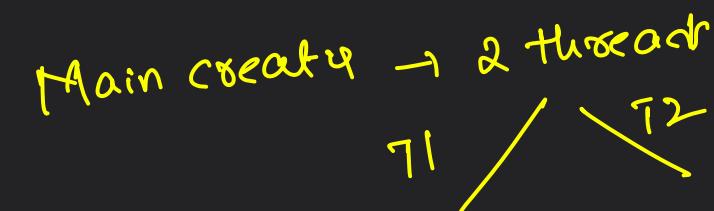
Now we need to overcome this problem, this can be done by **mutual exclusion (mutex)**



LOCKS (Mutex)

Let's understand working behind this
 $\text{Init} \rightarrow \text{Acquire} \rightarrow \text{Hold} \rightarrow \text{Release} \rightarrow \text{Destroy}$.

- 1) Init: $\text{pthread_mutex_init}(\&\text{lock}, \text{NULL})$
 Create mutex object & set the initial state to be unlocked



1) currently lock is free

- 2) T_1 calls $\text{pthread_mutex}(\&\text{lock})$
 - 3) T_1 gets lock
 - 4) T_1 does $++(\text{counter part})$
 - 5) T_1 calls unlock
 - 6) T_2 is in waiting state, while T_1 had the lock
 - 7) once T_1 is released T_2 gets the lock
 - 8) T_2 does $counter++$
 - 9) T_2 releases lock
- Lock is free again
- 500

$$\boxed{\text{Total exact } 500+500 = 1000 \text{ times}}$$

Semaphores: A variable which can be used as a shared variable to prevent race condition.

* Implemented through Wait() & Signal()

↓
Both are atomic in
nature

Wait() ↳ Gives the status of process to enter into critical section

* If Semaphore value is ≥ 1 then it can avail critical section

* Semaphore value will be decremented.

* If Semaphore val is '0', then indicates some process is in critical section. No other process is allowed

Signal(): It is an Increment operation

$$S = S + 1$$

Semaphore ↴

- ↳ Binary Semaphore (0, 1)
- ↳ Counting Semaphore (0, 1, ..., n)

```

1 #include <stdio.h>
2 #include <pthread.h>
3 #include <semaphore.h>
4 #include <unistd.h>
5
6 /*
7  * SIMPLE PROGRAM: two threads share one printer.
8  * Comments are written step-by-step (from main -> threads -> exit)
9  * and explain why each call is made and what the library does internally.
10 */
11
12 /* Global semaphore representing the printer.
13  * We treat it as a binary semaphore by initializing to 1 (free).
14  * sem_t holds: an integer counter + kernel-managed wait-queue (handled by
15  * sem_wait()).
16 */
17 sem_t printer;
18
19 /* Thread-function: represents a user submitting a print job.
20  * Steps inside this function are numbered and explained.
21 */
22 void* print_job(void* arg)
23 {
24     char* name = (char*) arg;
25
26     /* STEP A: Try to acquire the printer.
27      * sem_wait is an atomic library/kernel call:
28      * - If semaphore value > 0: sem_wait atomically decrements it (value--)
29      * and returns immediately.
30      * - If semaphore value == 0: sem_wait blocks this thread and the kernel
31      * puts it on the semaphore's wait-queue.
32      * Important: you do NOT write "printer--" yourself -- the decrement and
33      * blocking logic is done inside sem_wait.
34      * Reason: manual decrement would cause race conditions; sem_wait is
35      * implemented to be atomic and safe.
36      */
37     sem_wait(&printer);
38
39     /* STEP B: Now the thread has exclusive access to the printer (critical
40      * section).
41      * Only one thread can reach this point at a time because semaphore value
42      * was decremented to 0.
43      * We simulate the printing work with sleep(2).
44      */
45     printf("%s is using the printer...\n", name);
46     sleep(2); // simulate printing time (critical section).
47     printf("%s finished printing.\n", name);
48
49     /* STEP C: Release the printer.
50      * sem_post is also an atomic call:
51      * - it increments the semaphore value (value++).
52      * - If there are threads blocked on this semaphore, the kernel wakes up
53      * (unblocks) one of them.
54      * Reason: sem_post signals "resource available" and transfers control to
55      * one waiting thread if present.
56      * Note: because semaphores don't enforce ownership, any thread can call
57      * sem_post (unlike a mutex).
58      */
59     sem_post(&printer);
60
61     return NULL;
62 }
63
64 int main()
65 {
66     pthread_t t1, t2;
67
68     /* STEP 1: Initialize semaphore.
69      * sem_init(&printer, pshared, value)
70      * - &printer : address of semaphore object
71      * - pshared = 0 : semaphore is shared between threads of the same
72      * process (not between processes)
73      * - value = 1 : initial count (1 means printer is free). Initializing to
74      * 1 makes this behave like a binary semaphore.
75      * Reason: we want exactly one user to use the printer at a time, so start
76      * with 1 available "slot".
77      */
78     sem_init(&printer, 0, 1);
79
80     /* STEP 2: Create thread t1 (User1)
81      * pthread_create tells the OS to start a new thread that runs
82      * print_job("User1").
83      * Scheduling is up to the OS: thread may run immediately or later.
84      * Reason: we spawn two concurrent actors to demonstrate contention on the
85      * printer.
86      */
87     pthread_create(&t1, NULL, print_job, "User1");
88
89     /* STEP 3: Create thread t2 (User2)
90      * Same as above. Now we have two threads that might race to acquire the
91      * semaphore.
92      * Reason: to show that only one thread can proceed while the other blocks
93      * inside sem_wait.
94      */
95     pthread_create(&t2, NULL, print_job, "User2");
96
97     /* STEP 4: What happens now (runtime behavior - no explicit code)
98      * - Either t1 or t2 reaches sem_wait first (depends on scheduler).
99      * - First arrives sees value == 1 => sem_wait decrements to 0 => it
100      * proceeds into the critical section.
101      * - The other arrives and finds value == 0 => sem_wait blocks that thread
102      * and it sleeps until sem_post.
103      * Reason: this ensures mutual exclusion on the printer resource.
104      */
105
106     /* STEP 5: Wait for threads to finish
107      * pthread_join(t1, NULL) blocks the main thread until t1 finishes.
108      * pthread_join(t2, NULL) blocks until t2 finishes.
109      * Reason: if main returned earlier, program might exit while threads are
110      * still running (bad).
111      * Also we must ensure semaphore is not destroyed while a thread could
112      * still use it.
113      */
114     pthread_join(t1, NULL);
115     pthread_join(t2, NULL);
116
117     /* STEP 6: Cleanup
118      * sem_destroy releases any resources associated with the semaphore object.
119      * Reason: good practice to destroy what you initialize. Only do this after
120      * all threads are done.
121      */
122     sem_destroy(&printer);
123
124     return 0;
125 }

```

Binary-Semaphore Example Snippet

```

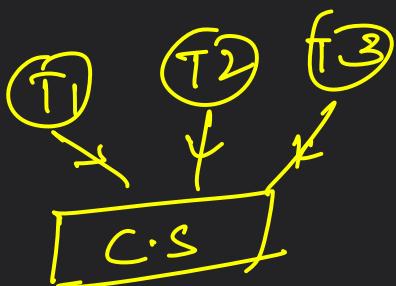
1 #include <stdio.h>
2 #include <pthread.h>
3 #include <semaphore.h>
4 #include <unistd.h>
5
6 /*
7 Example: 6 students share 3 library tables.
8 Counting semaphore initialized with 3 (for 3 tables).
9 We will trace how the semaphore value changes with each sem_wait / sem_post.
10
11 Compile: gcc -pthread semaphore_tables.c -o semaphore_tables
12 Run: ./semaphore_tables
13 */
14
15 #define NUM_TABLES 3
16 #define NUM_STUDENTS 6
17
18 sem_t tables; // counting semaphore
19
20 void* student(void* arg)
21 {
22     int id = *(int*)arg;
23
24     /* STEP A: Try to acquire a table
25      sem_wait:
26          - If tables > 0 → decrement tables-- and continue
27          - If tables == 0 → block until someone calls sem_post
28
29     Example value changes:
30         Start = 3
31         First student arrives → sem_wait → 3 → 2
32         Second student → sem_wait → 2 → 1
33         Third student → sem_wait → 1 → 0
34         Fourth student → finds 0 → BLOCKED
35
36     */
37     sem_wait(&tables);
38     printf("Student %d got a table (semaphore decremented).\n", id);
39
40     sleep(2); // simulate study time
41
42     /* STEP B: Release the table
43      sem_post:
44          - Increment tables++
45          - If any students are waiting, wake one of them
46
47     Example value changes:
48         When first student leaves → 0 → 1
49         Wakes up Student 4 → immediately sem_wait inside that student → 1 → 0
50
51     */
52     sem_post(&tables);
53
54     return NULL;
55 }
56
57 int main()
58 {
59     pthread_t tid[NUM_STUDENTS];
60     int ids[NUM_STUDENTS];
61
62     /* STEP 1: Initialize semaphore
63      sem_init(&tables, 0, 3)
64      Current value = 3 (3 tables free)
65
66     */
67     sem_init(&tables, 0, NUM_TABLES);
68
69     /* STEP 2: Create 6 student threads
70      Each will try sem_wait:
71          - First 3 succeed immediately (3→2→1→0)
72          - Remaining students block until tables free
73
74     */
75     for (int i = 0; i < NUM_STUDENTS; i++)
76     {
77         ids[i] = i + 1;
78         pthread_create(&tid[i], NULL, student, &ids[i]);
79     }
80
81     /* STEP 3: Join all threads
82      Main thread waits until all students finish.
83      Semaphore value will return to 3 at the end:
84          - Every sem_wait (--) is matched by a sem_post (++).
85
86     */
87     for (int i = 0; i < NUM_STUDENTS; i++)
88     {
89         pthread_join(tid[i], NULL);
90     }
91
92     /* STEP 4: Destroy semaphore (cleanup) */
93     sem_destroy(&tables);
94
95     return 0;
96 }

```

Counting Semaphore Example

Disadvantages of Locks

- ① Contention



T1 is under execution

- ① Assuming T1 has encountered an exception & it got terminated/dead, T2 & T3 will be in infinite wait state.

- ② Deadlock:



T1 has R1, but needs R2

T2 has R2, but needs R1

Each thread is waiting for other thread to unlock & deal.

Conclusion: To manage critical section problem, we are introducing locking mechanism.

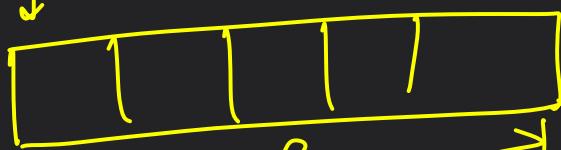
If locking mechanisms are not implemented correctly, it will lead to deadlock conditions.

Producers - Consumer problem

① Thread 1 → producer

② Thread 2 → consumer

producer



producer,
produce
data &
puts inside buffer

Buffer / critical
section

'n' slots

consumer

producer, produce data & fills in the empty slot
consumer, consumer data from the available slots

* Now i need to sync b/w P & C

* If n slots are full, we can't tell P to

fill the slots

(producer must not insert data
when buffer is full)

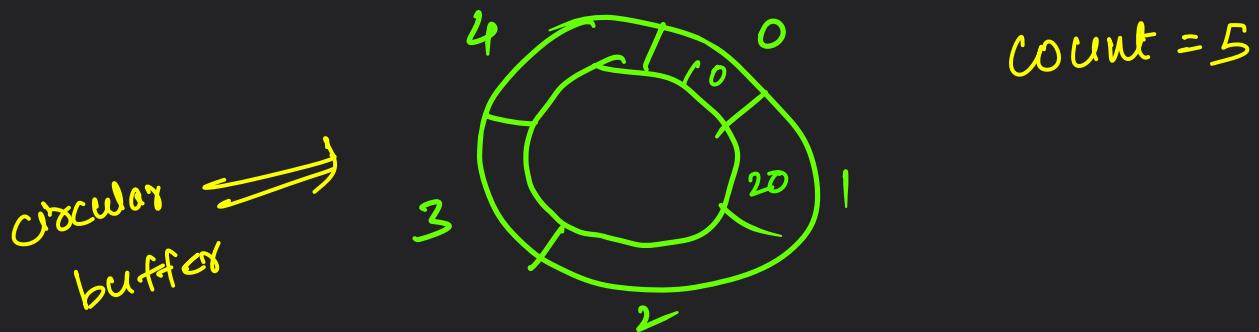
* Consumer must not pick / remove data when buffer
is empty.

Let solve this through Semaphore

① mutex m → Binary Semaphore,
used to acquire lock on
buffer

② counting Semaphore → initial val is 'n'
(empty)

③ Full \rightarrow Tracks filled slot
initial = 0



producer process \rightarrow In

consumer process \rightarrow Out

Code for producer problem

while(true)

{

 while(count == Buffer_size); \rightarrow 0 == 5 X

 buffer[in] = next-produced; \rightarrow buffer[0] = 10;

 in = (in + 1) % Buffer_size \rightarrow Suppose in = 4, next it should point to 0

 count++; \rightarrow count = 1

y

Last count is 5

5 == 5 ✓

[Infinite while loop]

Producer full (buffer full)

Iterations

count = 0, in = 0

0 == 5 X

buffer[0] = 10;

Suppose in = 4, next it should point to 0
 $3 \cdot 1 \cdot 5 = 0 \vee$

1
buffer[1] = 20;
in = 2

Now count = 2

Code for Consumer process

while (true)

```

d   while (count == 0); Buffer empty
    next_consumed = buffer[out];
    out = (out + 1) % Buffer-size;
    count--;
    // consume the item to be produced
  
```

y

$S = = 0 \times$

$out = 0;$

$\diamond next_consumed = 10;$

incr out

$C--;$

$U = = 0 \times$

$out = 1$

$next_cons = 20;$

incr out; // 2

$C--;$

Last iteration

$C = 0$

$O = = 0 \checkmark$

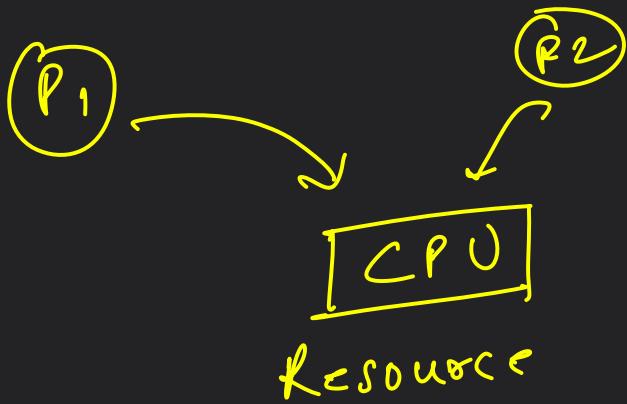
Infinite while

① producer problem buffer full
(producer must wait)

② consumer problem buffer empty
(consumer must wait)

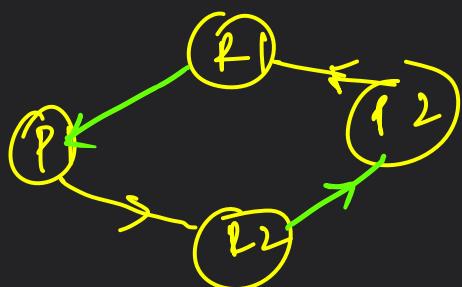
③ critical section (If producer & consumer try
to access the shared buffer
at same time)

Deadlocks



→ finite no. of resources are there

→ Multiple processes/thread



$P_1 \rightarrow R_1 \& R_2$
 $P_2 \rightarrow R_2 \& R_1$

Both are required

For acquiring resource, we use locking mechanism

$P_1 \rightarrow R_1$ lock

Now P_1 is waiting for R_2

$P_2 \rightarrow R_2$ lock

Now P_2 is waiting for R_1

* This waiting mechanism can lead to infinite wait time.

* Now during this situation, the system is in deadlock.

How a process/thread utilize a resource?

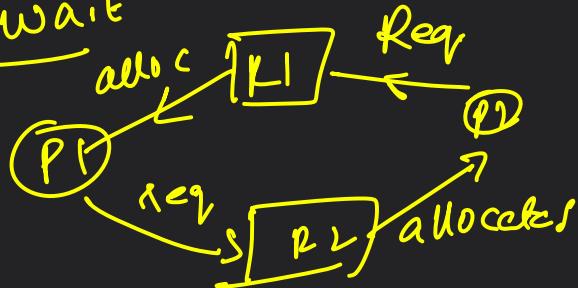
- ① Request
- ② Use
- ③ Release

Necessary condition for deadlock

① Mutual exclusion

1 Resource \rightarrow 1 process / 1 thread

② Hold & wait



P1 is holding R1

also

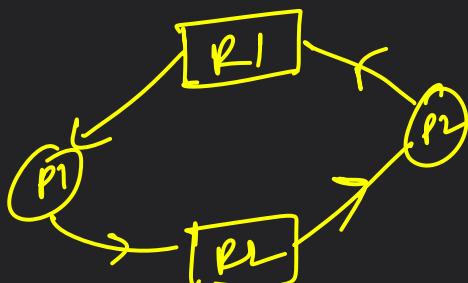
P1 is waiting for R

③ No preemption:



P1 \rightarrow Execution complete

④ Circular wait



Resource allocation graph:

① Vertices \rightarrow ① Process Vertex

② Resource Vertex

③ Edges \rightarrow ① Assign Edge
② Request



Methods to handle Deadlock:

① Prevent mutual exclusion

Eg Replicate/Duplicate the resource
Multiple UART ports, tasks can use different ones.

② Make the resource read only

Eg If no state change happens multiple process can read at the same time
(files opened in read only mode)

Deadlock Avoidance

① Current state of System

- ↳ No. of processes
- ↳ Need of Resources of each process
- ↳ currently allocated amount of R
- ↳ Max amount of each Resource

OS scheduler makes decisions dynamically at run-time
ensure that system never reaches unsafe state

↳ less safe to update ref to
no. of resources required to complete

Total Resource = 16
Current available = 3

P1: Max need = 7
Current alloc = 4
Need 3 to complete

P2: Max need = 4
Current alloc = 2
Need 2 more to complete

Safe-state:

Give P2, 2 resources (current available
 $3 - 2 = 1$)

P2 finishes and releases
So total available ($4 + 1 = 5$)

Give P1 → 3 [So total available
 $5 - 3 = 2$]

and $P1 = 4 + 3 = 7$

P1 completed

Case 2: Unsafe

P1 requesting 3, and we give it directly 3 to it

But P2 is requesting 2 to complete

Now either P1 or P2 will finish
Deadlock

Deadlock detection:

Assuming our system is not having the capability of doing prevention or avoidance, so during these cases we manually implement a deadlock detection mechanism (OS uses resource allocation graph to check the cycles)

Deadlock Recovery

① Process Termination:

In System we are having 10 processes 3 are suffering with deadlock

Kill -9 3 processes

CPU time wastage, Memory wastage
I/O device dependent on this
also wastage-

② Abort 1 process

check deadlock detection algorithm

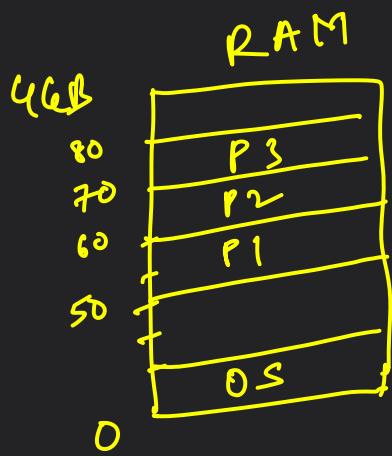
Repeat

overhead of detection algorithm

③ Resource-preemption : Forcefully release the resources

- a) Selecting a scheduler
- b) Rollback (execute from scratch)
- c) Starvation (Resource starvation)

Memory Management



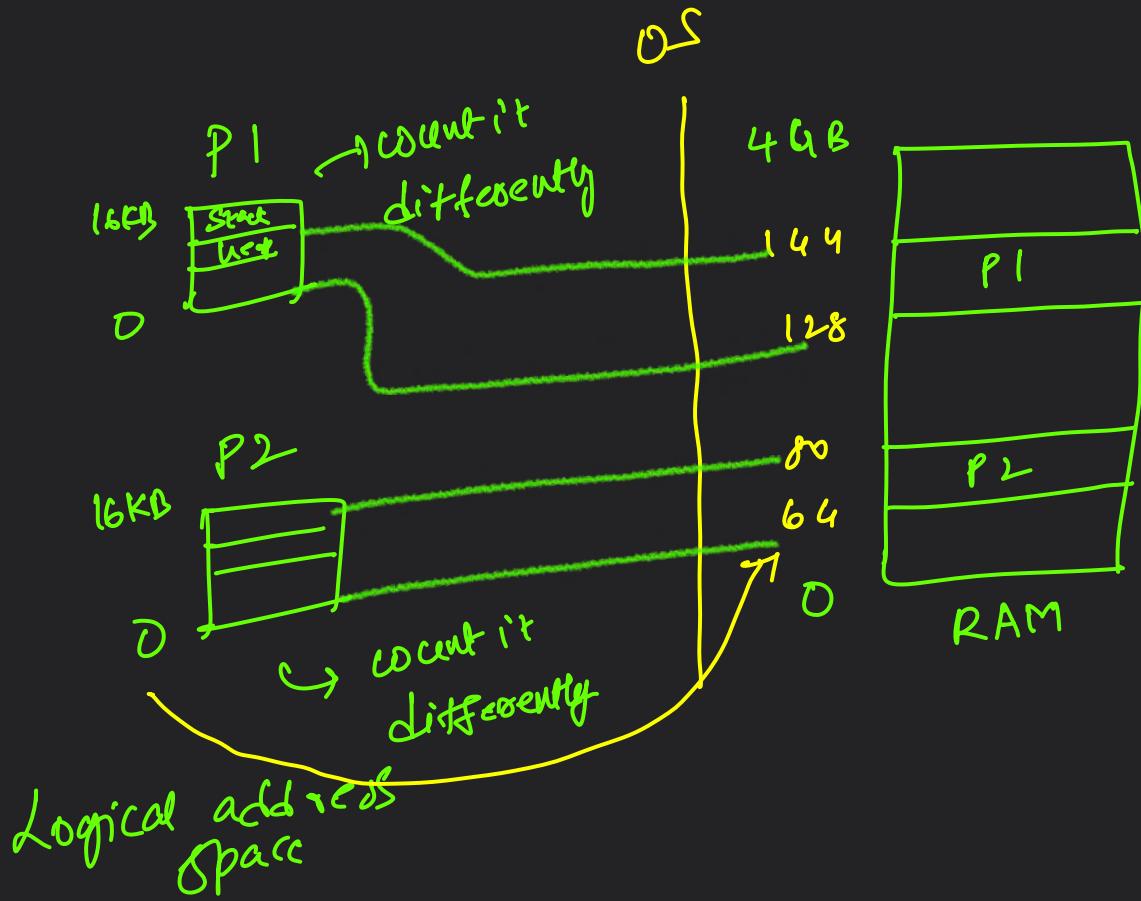
P1 Address Space
50 - 60

Suppose P1 is trying to do
P1 address + 5

$$60 + 5 = 65$$

So P1 is trying to access P2

(This happens when OS is not having protection & memory isolation)



OS → Base & offset
 ↓ ↓
 128 16 (P1)

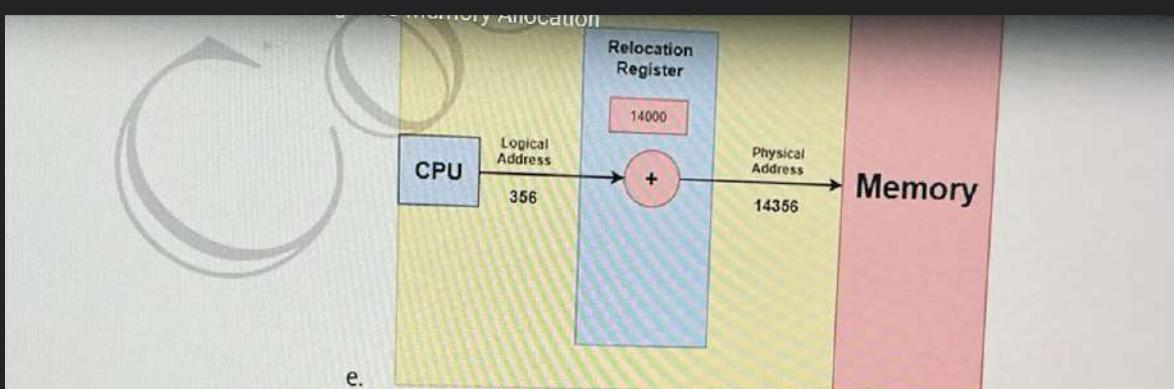
OS does the mapping of logical to physical address
 (Address Translation)

Logical addresses are generated by CPU

* OS & MMU computes the physical address

* MMU is a hardware device.

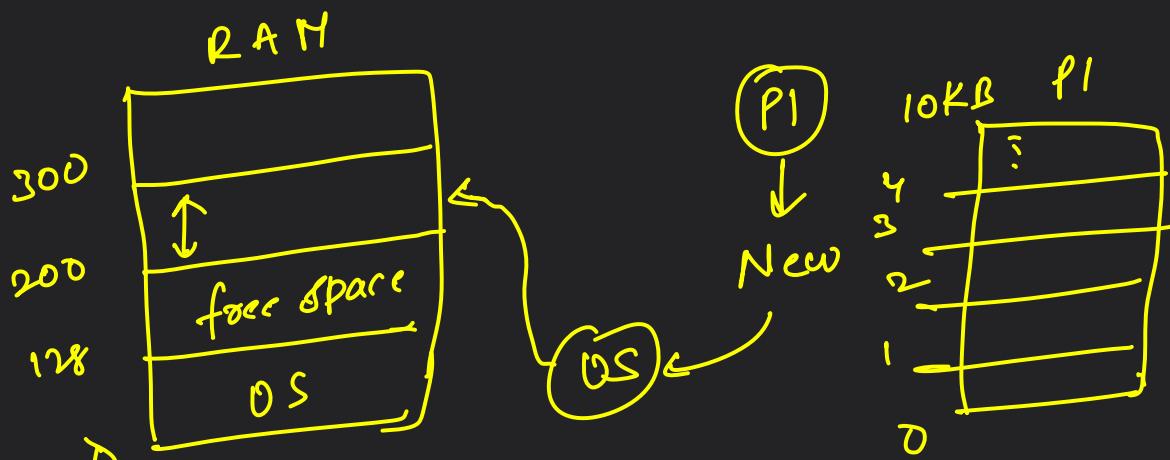
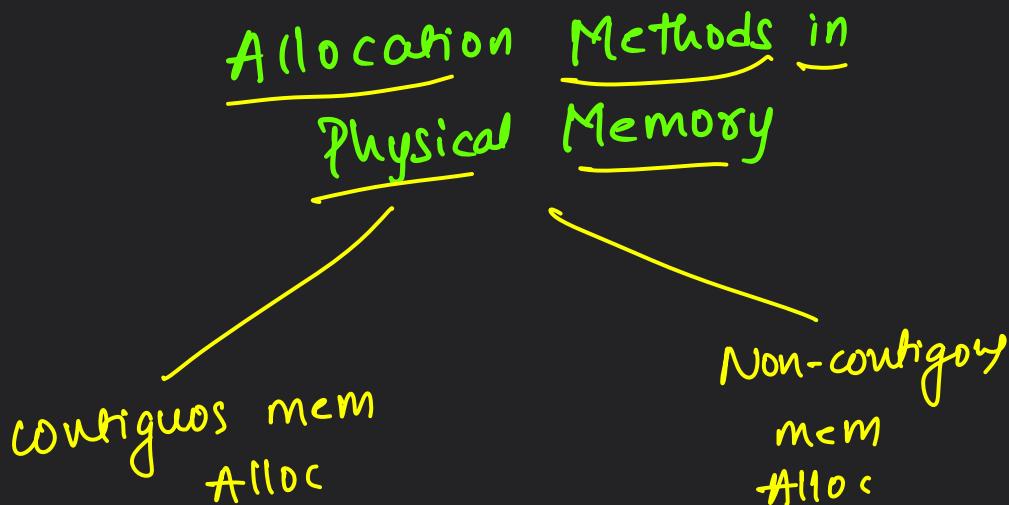
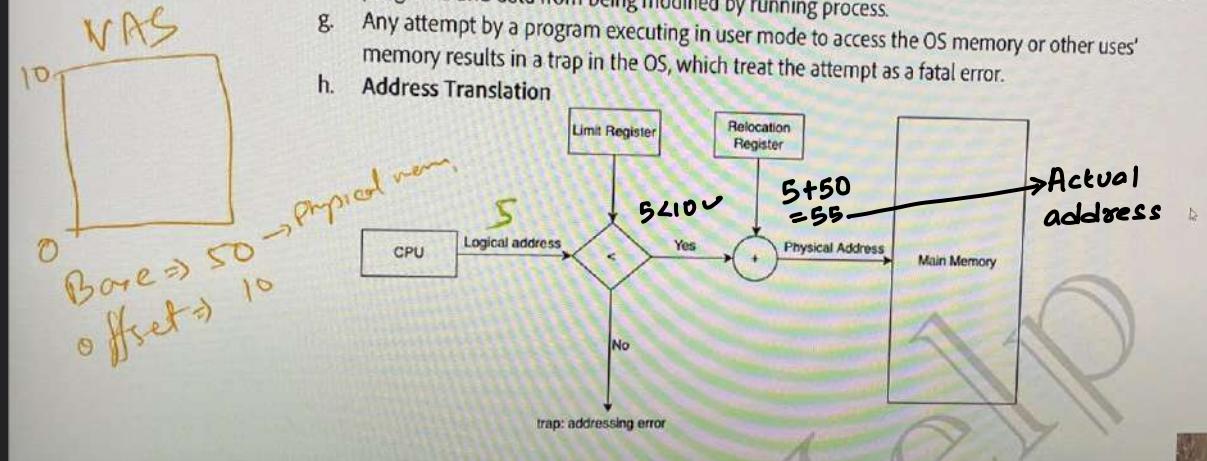
* Program needs physical address in order to execute.



4. How OS manages the isolation and protect? (Memory Mapping and Protection)

- a. OS provides this Virtual Address Space (VAS) concept.
- b. To separate memory space, we need the ability to determine the range of legal addresses that the process may access and to ensure that the process can access only these legal addresses.
- c. The relocation register contains value of smallest physical address (Base address [R]); the limit register contains the range of logical addresses (e.g., relocation = 100040 & limit = 74600).
- d. Each logical address must be less than the limit register.

- e. MMU maps the logical address dynamically by adding the value in the relocation register.
- f. When CPU scheduler selects a process for execution, the dispatcher loads the relocation and limit registers with the correct values as part of the context switch. Since every address generated by the CPU (Logical address) is checked against these registers, we can protect both OS and other users' programs and data from being modified by running process.
- g. Any attempt by a program executing in user mode to access the OS memory or other users' memory results in a trap in the OS, which treat the attempt as a fatal error.
- h. Address Translation



continuously (contiguously)
Put P1 to memory (RAM)

① Fixed partitioning of RAM

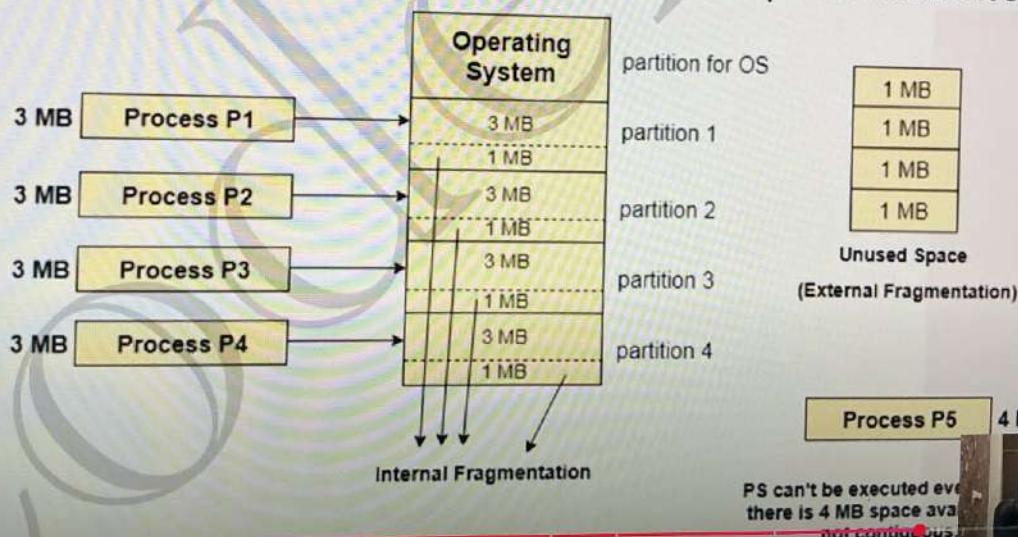
Q

variable size



Fixed Partitioning

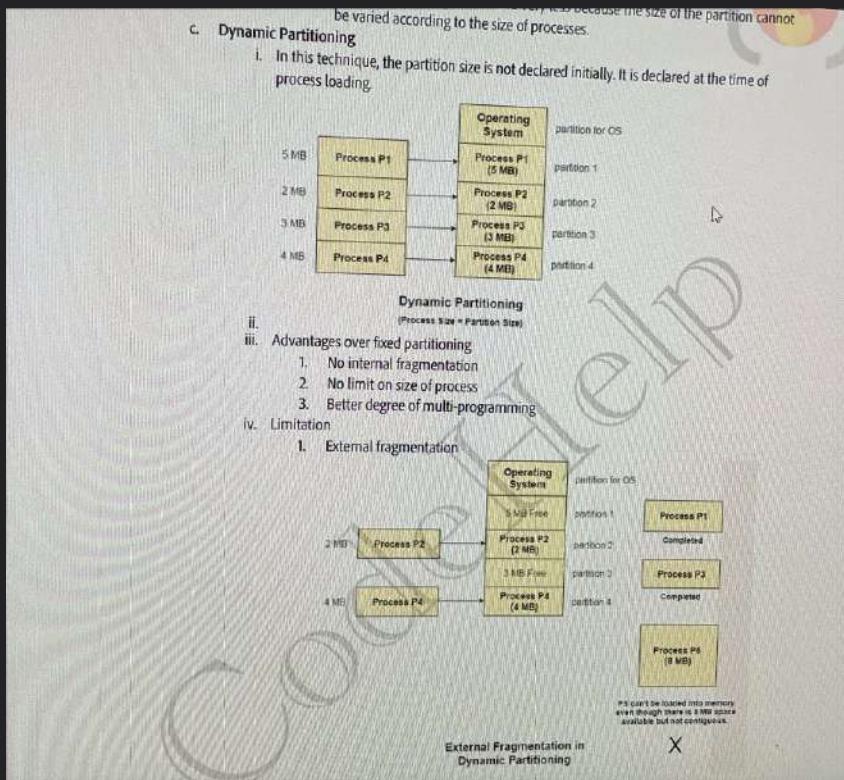
- The main memory is divided into partitions of equal or different sizes.



iii. Limitations:

- Internal Fragmentation:** if the size of the process is lesser than the total size of the partition then some size of the partition gets wasted and remain unused. This is wastage of the memory and called internal fragmentation.
- External Fragmentation:** The total unused space of various partitions cannot be used to load the processes even though there is space available but not in the contiguous form.
- Limitation on process size:** If the process size is larger than the size of maximum sized partition then that process cannot be loaded into the memory. Therefore, a limitation can be imposed on the process size that is it cannot be larger than the size of the largest partition.

Next we will be discussing on dynamic partitioning.



P1 & P3 are completed

$$P_1 = 5 \text{ MB}$$

$$P_3 = 3 \text{ MB}$$

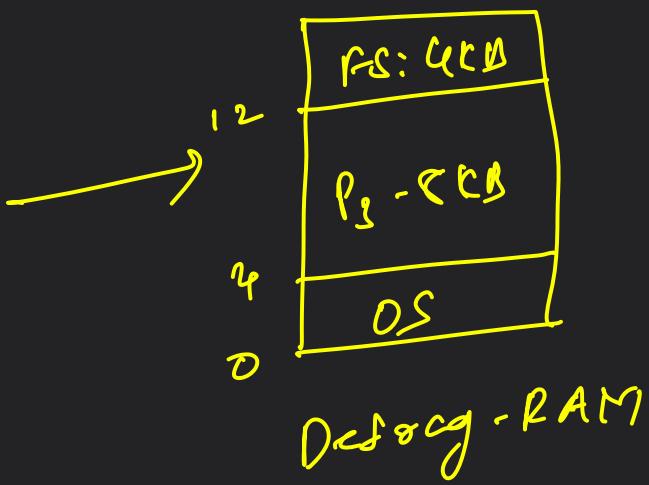
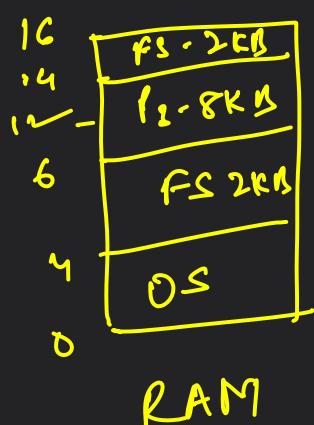
We have 8 MB free space

so now P5 comes with 8 MB, it can't unless the above 8 MB

* For free space memory mgmt we are using a free list

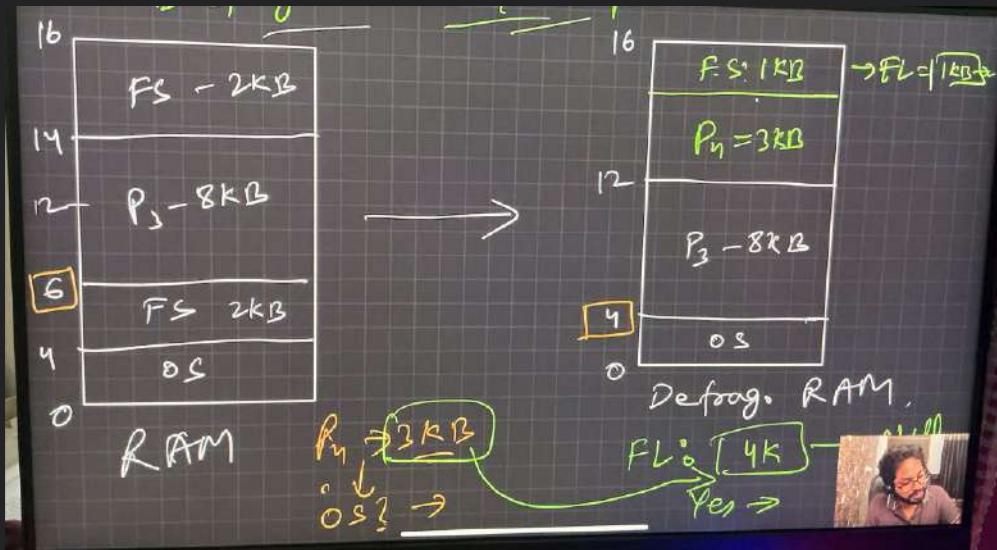
* It is a type of linked list

Defragmentation / compaction:



Re-arranges scattered files/blocks into continuous memory locations (once files are deleted, free space get scattered into

Small chunk

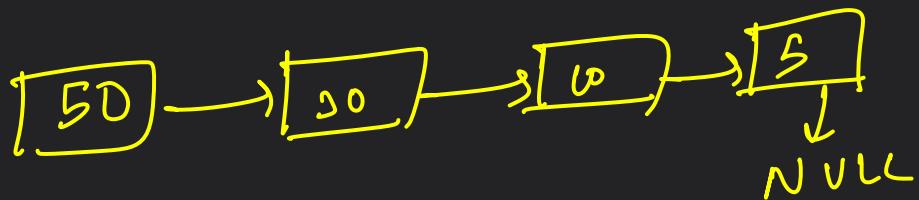


Take the chunks, combine them and do contiguous location

* Time consuming (Defragmentation)

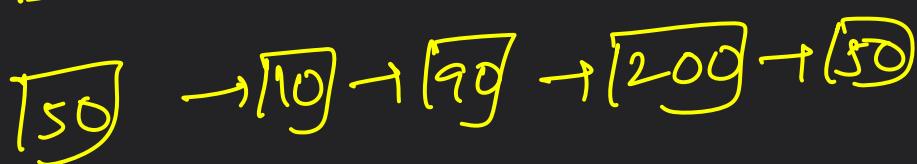
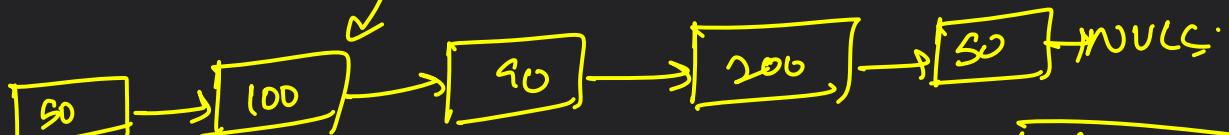
Just like RAM, Disk can also be defragmented.

Free-space called as holes represented through
Linked List



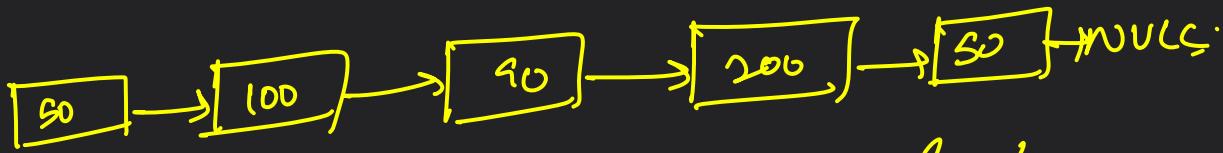
For 20KB how to fit in free list?

① First fit 90KB (Big enough first)



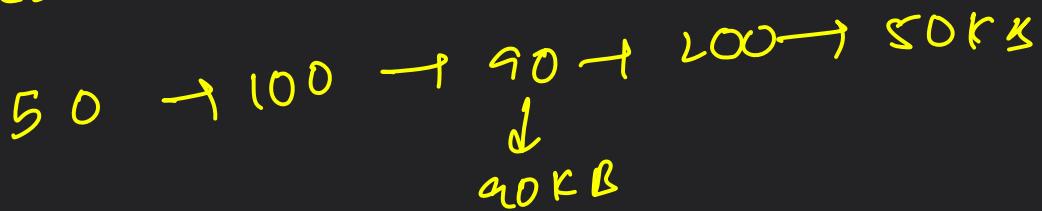
fast
simple

② Next fit

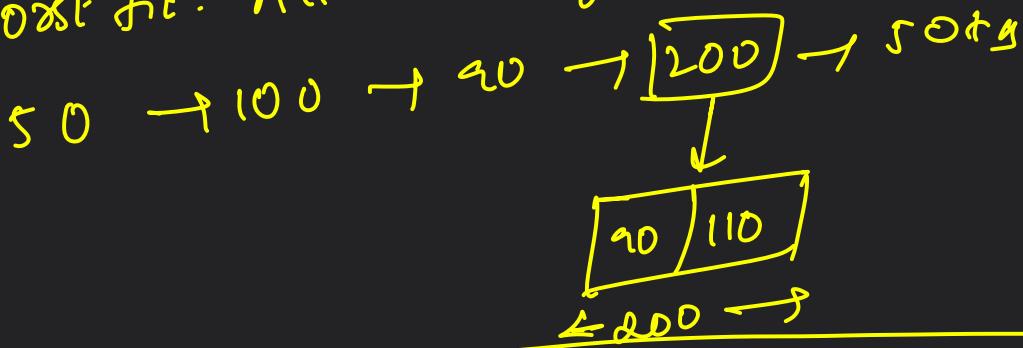


$P_2 \rightarrow 100\text{KB}$ (Previous saved next loc)

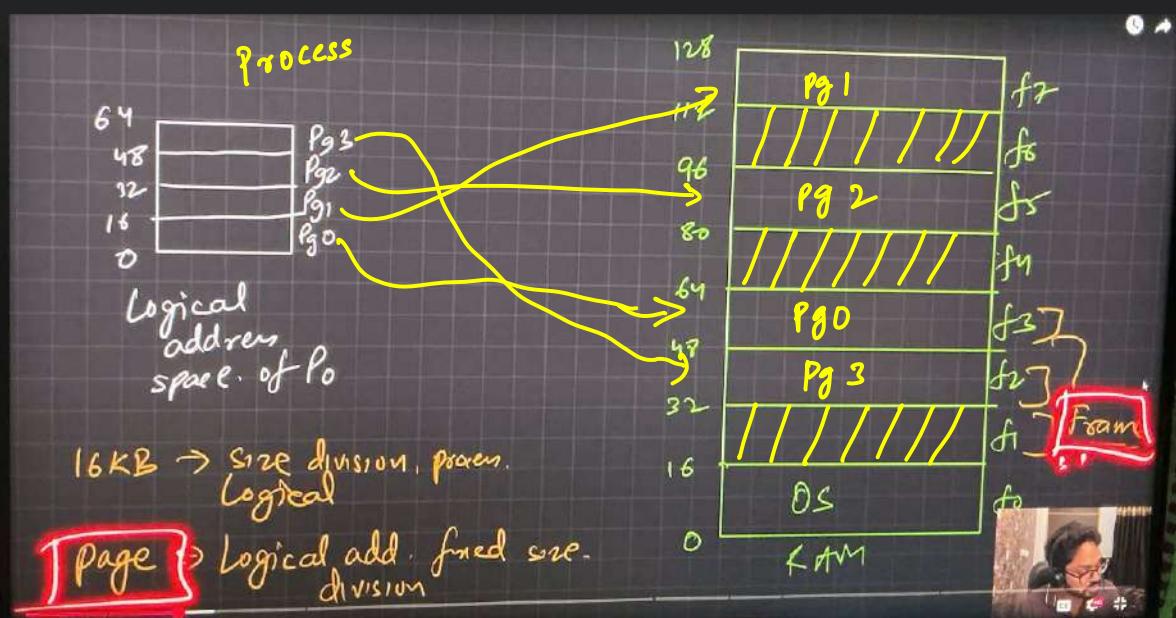
③ Best fit (smallest hole that is big enough)



④ Worst fit: Allocate largest hole



Paging (Non-contiguous)



Page size = frame size

Page Table

Logical page no

0 0	Pg 0
0 1	1
1 0	2
1 1	3

Physical frame
frame 3

7
5
2

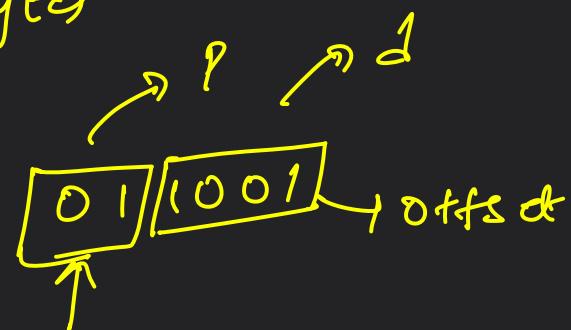
L.A \rightarrow Pg 0 \rightarrow 64 bytes

$$2^6 = 64$$

6 bits $\Rightarrow 2^5$

\downarrow

Logical address as
Page no.



Page no + Base \rightarrow 16 + 9 $\Rightarrow 25$
(offset)

Paging in Operating Systems - Memory Management

Paging Overview -

- In computer operating systems, paging is a memory management scheme by which a computer stores and retrieves data from secondary storage for use in main memory.
- In this scheme, the operating system retrieves data from secondary storage in same-size blocks called pages.
- Paging is an important part of virtual memory implementations in modern operating systems, using secondary storage to let programs exceed the size of available physical memory.
- Non-contiguous memory allocation
- Helps prevent external fragmentation
- Logical address space is divided into equal size pages
- Physical address space is divided into equal size frames
- Page Size = Frame Size

>> Logical Address or Virtual Address (represented in bits) - An address generated by the CPU

>> Physical Address (represented in bits) - An address actually available on memory unit

>> The mapping from virtual to physical address is done by the memory management unit (MMU) which is a hardware device and this mapping is known as paging technique.

